

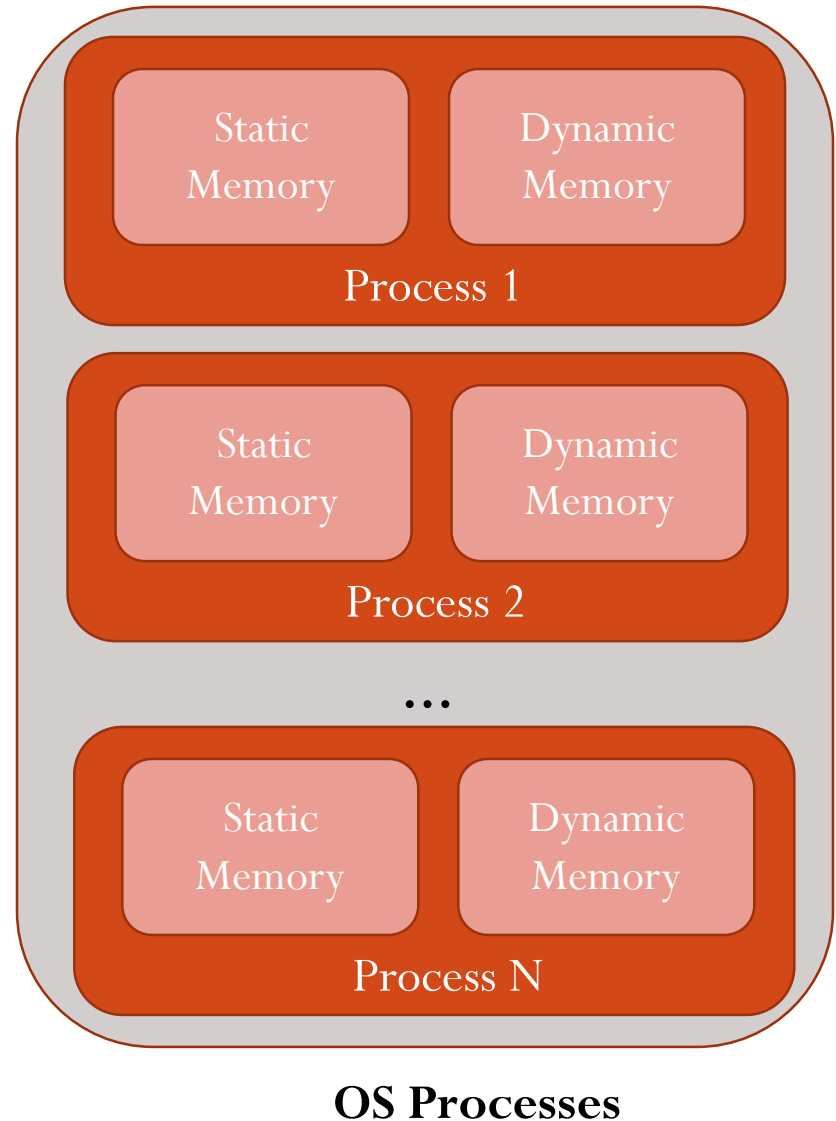
Dynamic Memory

Chapter 4

Pages 153-167

Shallow Dive into Memory

- There are two memory pools for programs running in a modern OS
- Static Memory Pool:
 - Referred to as the stack
- Dynamic Memory Pool:
 - Referred to as the heap



The Stack (Static Memory)

- Stack Memory:
 - On each encounter of a function or new scope declared variables are placed on top of the stack.
 - On exiting the scope variables are popped off the stack
 - The top of the stack is tracked by the stack pointer which is a simple counter
 - Variables can only be popped from the top of the stack, not at any other location

The Stack (Static Memory)

- Lets investigate the stack allocation of the code below:
- Here we start with an empty stack

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}
```

```
void func1(int x){  
    int y;  
    func2();  
}
```

```
void func2(){  
    int z[2];  
}
```

stkPtr →

Mem Addr	Memory
0x1000	
0x1004	
0x1008	
0x100c	
0x1010	
0x1014	
0x1018	

The Stack (Static Memory)

- int a, b are allocated by placing them on the stack within a function frame
- The stack pointer decreases accordingly

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}
```

```
void func1(int x){  
    int y;  
    func2();  
}
```

```
void func2(){  
    int z[2];  
}
```

stkPtr →

Mem Addr	Memory
0x1000	a = 2
0x1004	b
0x1008	
0x100c	
0x1010	
0x1014	
0x1018	

} main func
frame

The Stack (Static Memory)

- Here we enter func1() through a call
- Parameters and declared variables of func1() are added to the stack

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```

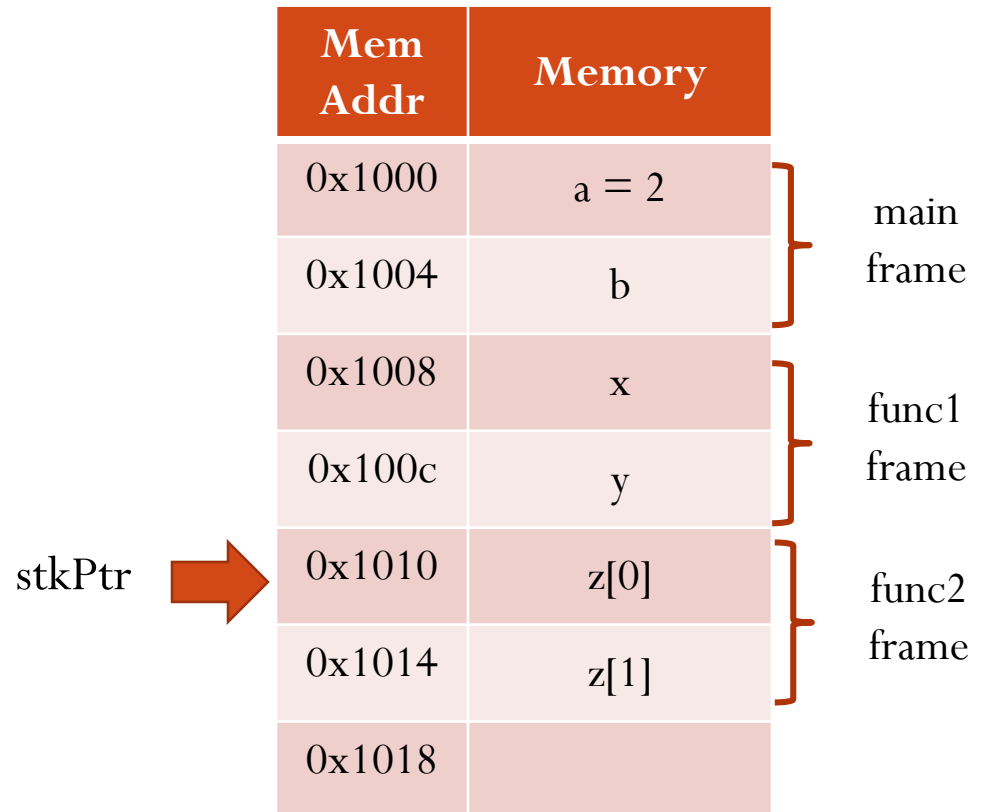
stkPtr →

Mem Addr	Memory	
0x1000	a = 2	main frame
0x1004	b	
0x1008	x	func1 frame
0x100c	y	
0x1010		
0x1014		
0x1018		

The Stack (Static Memory)

- Now we enter func2() through func1()
- Parameters and declared variables of func2() are added to the stack

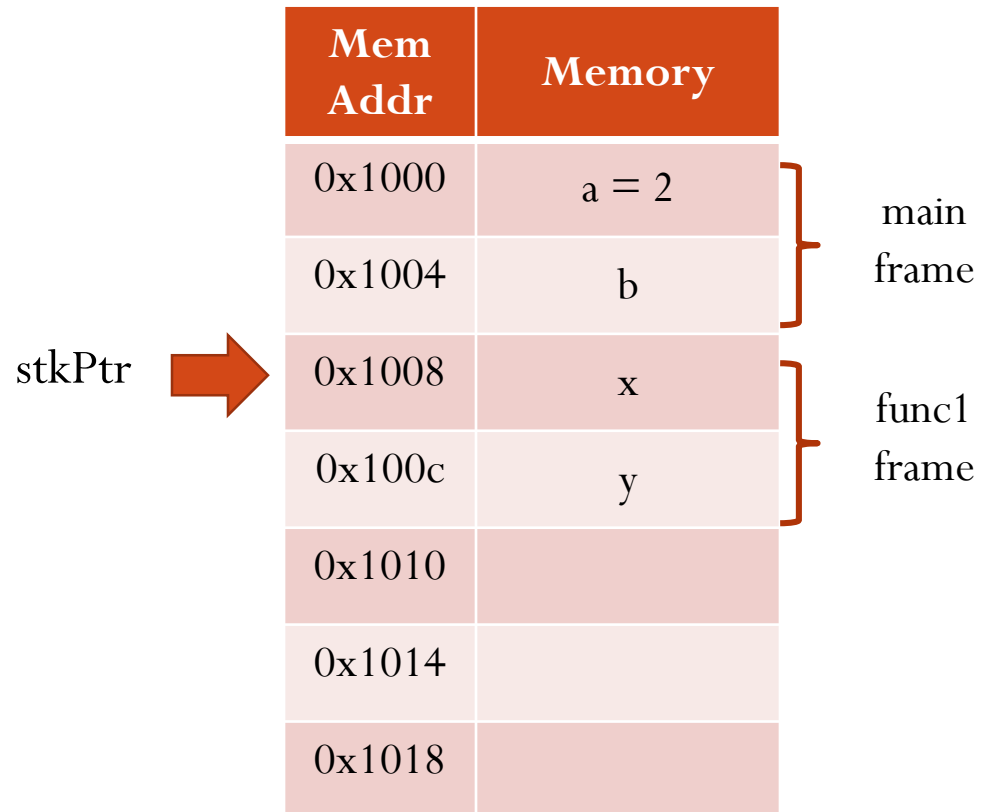
```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```



The Stack (Static Memory)

- Now func2() is left returning to func1()
- Variables on the stack are popped accordingly increasing our stkPtr

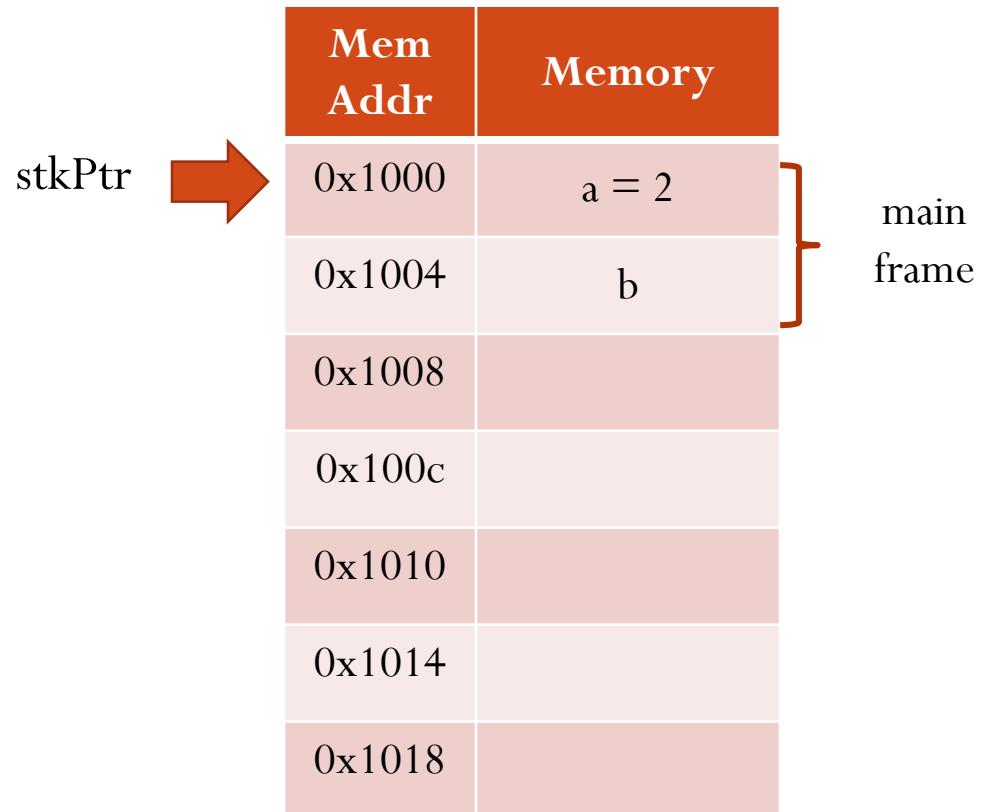
```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```



The Stack (Static Memory)

- Now func1() is left returning to main()
- This bring our stack to our original variables

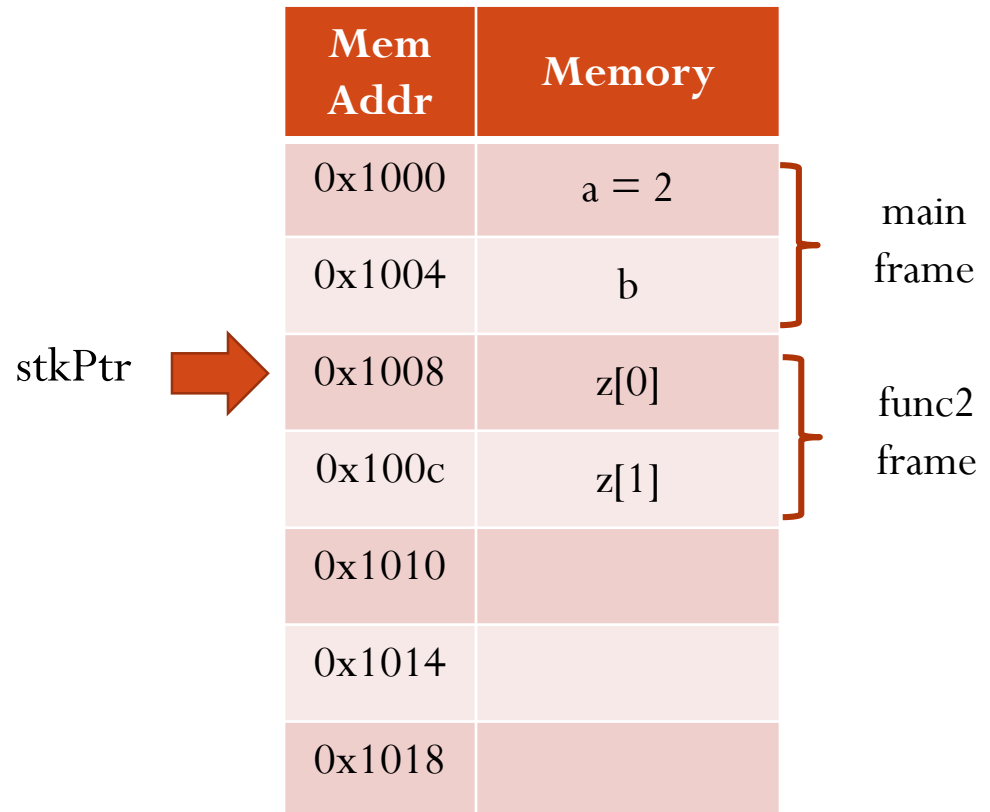
```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```



The Stack (Static Memory)

- The program enters func2() again through main()
- Add variables back in the stack

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```



The Stack (Static Memory)

- Program exits func2()
- Data from func2() is removed from the stack

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```

stkPtr →

Mem Addr	Memory	}	main frame
0x1000	a = 2		
0x1004	b		
0x1008			
0x100c			
0x1010			
0x1014			
0x1018			

The Stack (Static Memory)

- Program exits main()
- Main data is removed leaving the stack empty

```
int main (){  
    int a = 2, b;  
    func1(a);  
    func2();  
    return 0;  
}  
  
void func1(int x){  
    int y;  
    func2();  
}  
  
void func2(){  
    int z[2];  
}
```

stkPtr →

Mem Addr	Memory
0x1000	
0x1004	
0x1008	
0x100c	
0x1010	
0x1014	
0x1018	

Some Fundamentals

- As functions are entered the variables are placed on the stack, the stack pointer decreases
- As functions are exited the variables are removed from the stack and the stack pointer increases
- **Variable locations are hardcoded into the assembly based on its relative location with respect to the stack pointer**
 - **Fixed at compile time!!!**
- What are the implications of this?

Implications of referring to stkPtr

- Lets jump to the first call of func1(2)
- What is the address of z?
 - $\text{stkPtr} + 8$

```
int main (){  
    int a = 2;  
    func1(a);  
    a = 3;  
    func1(a);  
    return 0;  
}  
  
void func1(int size){  
    int y[size];  
    int z;  
}
```

stkPtr →

Mem Addr	Memory	
0x1000	a = 2	main frame
0x1004	y[0]	
0x1008	y[1]	func1 frame
0x100c	z	
0x1010		
0x1014		
0x1018		

Implications of referring to stkPtr

- Now the second call of func1(3)
- What is the address of z?
 - $\text{stkPtr} + 12$

```
int main (){  
    int a = 2;  
    func1(a);  
    a = 3;  
    func1(a);  
    return 0;  
}  
  
void func1(int size){  
    int z;  
    int y[size];  
}
```

stkPtr



Mem Addr	Memory	
0x1000	a = 2	main frame
0x1004	y[0]	
0x1008	y[1]	func1 frame
0x100c	y[2]	
0x1010	z	
0x1014		
0x1018		

Implications of referring to `stkPtr`

- In the two examples the address of `z` was:
 - `strPtr + 8`
 - `strPtr + 12`
- If the size of the array is determined at runtime is it possible to hardcode the addresses of variables based on its relative position to the stack pointer?
 - No, because the relative position will also change at runtime
- Implication:
 - The size of the array must be known at compile time if it is to be placed on the stack

The Heap (Dynamic Memory)

- The heap is a less rigid memory storage structure
- Unlike the stack, individual memory allocations are **not** guaranteed to be in sequential order
- Bookkeeping is performed by a heap memory manager
 - A search must be performed to find a memory location large enough to store the information
- The heap offers greater flexibility since the size of the memory allocations only need to be known at runtime

A Heap of Terms


- Allocation:
 - A reservation of one or several consecutive memory locations in the heap for data storage
- Deallocation:
 - Freeing of a previous memory allocation (reservation) to allow for future reservations of the same location by different data
- Memory Leak:
 - An allocation of memory that is never deallocated throughout the life of the program
 - For every allocation, there should always be a deallocation!!!

The Heap

- Up until now we have exclusively used the stack with some exceptions:
 - C++ strings
 - ever wondered why the library could get away with runtime sizing of character arrays but you couldn't?
 - Cout, Cin, fstreams
- Allocations to the heap can be performed using the following:
 - C++: keyword *new*
 - C: function `malloc()`
- Deallocations to the heap can be performed using:
 - C++: keyword *delete*
 - C: function `free()`

Heap Visualization

- Given the unstructured nature, the heap can be visualized as a cloud
- While there are complex algorithms to allocate memory we shall assume near random assignment of allocating memory



Using Dynamic Memory

```
int main ()
{
    int *num;
    num = new int;

    *num = 5;
    cout << *num << endl;
    delete num;

    return 0;
}
```

- To the right is an example of a dynamic allocation of an integer, and deallocation of an integer
- Notice we are working with a pointer
- Lets step through the process

Using Dynamic Memory

```
int *num;
```

- Start with a pointer on the stack and an empty heap

Mem Addr	Variable
0x1004	
0x1000	num = ?
The Stack	


The Heap

Using Dynamic Memory

```
int *num;  
num = new int;
```

- The keyword *new* returns the memory address of the allocated memory

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	




Using Dynamic Memory

```
int *num;  
num = new int;
```

- Thus we have a pointer on the stack pointing to our reserved memory

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	




Using Dynamic Memory

```
int *num;  
num = new int;  
*num = 5;
```

- Now lets put some data in there using a dereference operator

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	




Using Dynamic Memory

```
int *num;  
num = new int;  
*num = 5;  
cout << *num << endl;
```

- The data can be fetched using the deref op.
- 5 will be printed to console

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	




Using Dynamic Memory

```
int *num;  
num = new int;  
*num = 5;  
cout << *num << endl;  
delete num;
```

- Finally the delete frees the memory in our heap
- Notice our pointer still points to 0x411C though

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	




The Heap

Using Dynamic Memory

```
int *num;  
num = new int;  
*num = 5;  
cout << *num << endl;  
delete num;
```

- Once the memory is freed, it would be a terrible idea to continue to try and use it

Mem Addr	Variable
0x1004	
0x1000	num = 0x411C
The Stack	



The Heap

Dynamic Memory Assignment


- The new operator returns a pointer to the allocated memory
- Steps for Allocation
 1. Declare pointer of data type of allocated memory
 - `float *x;`
 2. Allocate memory using new keyword and the data type to allocate
 - `x = new float;`

Allocations

```
double *num1;  
int *num2;  
num1 = new double;  
num2 = new int;  
*num1 = 5.134;  
*num2 = 10;
```

- Multiple allocations

Mem Addr	Variable
0x1004	num2 = 0x800A
0x1000	num1 = 0x7224
The Stack	




Deallocations

```
double *num1;  
int *num2;  
num1 = new double;  
num2 = new int;  
*num1 = 5.134;  
*num2 = 10;  
delete num1;  
delete num2;
```

- Multiple deallocations

Mem Addr	Variable
0x1004	num2 = 0x800A
0x1000	num1 = 0x7224
The Stack	



The Heap

Dynamic Allocation of Arrays

- An array can be allocated on the heap by requesting a larger allocation
- This can be achieved by altering the syntax of the allocation

```
int *arr;  
arr = new int[10];
```

- More generically

```
datatype *arr;  
arr = new datatype[#_of_elements];
```

- Since the heap is dynamic, the number of elements can be a variable

Array Allocation

```
float *arr;  
arr = new float[4];
```

- Here we have pointer to an allocation of 4 consecutive floats in our heap

Mem Addr	Variable
0x1004	
0x1000	arr = 0x7224
The Stack	

0x7230	
0x722C	
0x7228	
0x7224	

The Heap

Dynamic Array Assignment

```
float *arr;  
arr = new float[4];  
arr[0] = 1.337;
```

- Once the allocated, assigning values into a dynamic array is identical to that of a static

Mem Addr	Variable
0x1004	
0x1000	arr = 0x7224
The Stack	

0x7230	
0x722C	
0x7228	
0x7224	1.337

The Heap

Dynamic Array Access

```
float *arr;  
arr = new float[4];  
arr[0] = 1.337;  
cout << arr[0] << endl;
```

- Again, fetching from a dynamic array does not differ to that of a static
- Here our code will print 1.337

Mem Addr	Variable
0x1004	
0x1000	arr = 0x7224
The Stack	

0x7230	
0x722C	
0x7228	
0x7224	1.337


The Heap

Deallocation of Dynamic Array

```
float *arr;  
arr = new float[4];  
arr[0] = 1.337;  
cout << arr[0] << endl;  
delete [] arr;
```

- Empty brackets are used to indicate a deallocation of an array
- Brackets must be placed before the array name
- Do not place a size inside the brackets

Mem Addr	Variable
0x1004	
0x1000	arr = 0x7224
The Stack	



The Heap


Dynamic Array with Variable Size

```
int size;  
double *grades;
```

```
cin >> size;  
grades = new double[size];  
...  
delete [] grades;
```

- Here we have the user select the size of the array
- Allocation is based on the variable size

Mem Addr	Variable
0x1004	arr = 0x8000
0x1000	size
The Stack	



...	
0x8008	
0x8004	
0x8000	

The Heap


Dynamic Array with Variable Size

```
int size;  
double *grades;
```

```
cin >> size;  
grades = new double[size];  
...  
delete [] grades;
```

- Notice the allocation can occur at any time
- This time it must at least be after size is initialized to a value

Mem Addr	Variable
0x1004	arr = 0x8000
0x1000	size
The Stack	



...	
0x8008	
0x8004	
0x8000	

The Heap

Example 15.1

- Create a program which calculates the mean and variance of a set of numbers
- Mean and Variance are given by:

$$Mean(x) = \frac{1}{n} \sum_{i=1}^n x_i$$

$$Var(x) = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$$

- Prompt the user for how many numbers they wish to enter
 - Create an array that will accommodate the requested numbers
- Fill the array through user input
- Print the mean, and variance


Reusing Pointers

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
delete [] arr;
```

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Allocate an array and place number 5 into it

Mem Addr	Variable
0x1004	
0x1000	arr = 0x8000
The Stack	



0x800C	
0x8008	
0x8004	
0x8000	5

The Heap


Reusing Pointers

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
delete [] arr;
```

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Deallocate the array after we are finished

Mem Addr	Variable
0x1004	
0x1000	arr = 0x8000
The Stack	



The Heap


Reusing Pointers

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
delete [] arr;
```

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Pointer *arr* is reused to create another array of a different size
- Notice *arr* points to a new location now

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6544
The Stack	



0x6554	3
0x6550	
0x654C	
0x6548	
0x6544	

The Heap


Reusing Pointers

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
delete [] arr;
```

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Pointer *arr* is reused to create another array of a different size
- Notice *arr* points to a new location now

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6544
The Stack	



0x6554	3
0x6550	
0x654C	
0x6548	
0x6544	

The Heap


Reusing Pointers

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
delete [] arr;
```

- Finally we deallocate the array

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6544
The Stack	




The Heap

Memory Leak

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
  
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Here we modify the code to forget a deallocation

Mem Addr	Variable
0x1004	
0x1000	arr = 0x8000
The Stack	



0x800C	
0x8008	
0x8004	
0x8000	5

The Heap


Memory Leak

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...
```

```
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- Here we perform another allocation but our previous allocation still exists

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6554
The Stack	



0x6554	3
0x6550	
0x654C	
0x6548	
0x6544	

0x800C	
0x8008	
0x8004	
0x8000	5


The Heap

Memory Leak

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- It is possible to delete the newest allocation but we have lost the pointer to our older one
- For the rest of the process's life the memory at 0x8000 will be reserved and unusable

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6554
The Stack	



0x800C	
0x8008	
0x8004	
0x8000	5


The Heap

Memory Leak

```
int *arr;  
arr = new int[4];  
arr[0] = 5;  
...  
arr = new int[5];  
arr[4] = 3;  
...  
delete [] arr;
```

- This is an example of a memory leak
- Rule of Allocations:
 - For every allocation there should be a deallocation

Mem Addr	Variable
0x1004	
0x1000	arr = 0x6554
The Stack	



0x800C	
0x8008	
0x8004	
0x8000	5


The Heap

How bad can a memory leak be?

- Consider the following:

```
int *arr, i;  
  
for(i = 0; i < 10000; i++)  
    arr = new int[1000];  
delete [] arr;
```

Mem Addr	Variable
0x1004	
0x1000	arr = ...
The Stack	



...	
0x####	

...	
0x####	

...	
0x####	

...

...	
0x####	

The Heap

How bad can a memory leak be?

- Consider the following:

```
int *arr, i;  
  
for(i = 0; i < 10000; i++)  
    arr = new int[1000];  
delete [] arr;
```

- How much memory is lost?
 - int is 32 bytes, each array is then $32 * 1000$, 32000 bytes
 - 10000 – 1 arrays were allocated without a deallocation
 - Therefore $32000 * 99999 = 319968000$ bytes
 - ≈ 320 MB

Memory Leaks

- While the example given may seem far fetched, often in practice memory allocation is performed in loops
- Such loops maybe executed millions of times
- Single instances of a memory leak are typically not an issue, however many repetitions of a single memory leak can spell disaster for your program
- Your program may run out of memory or worse effect the stability of your operating system

Memory Leaks

- Memory leaks are very hard to detect manually for the following reasons
 - They do not effect the functionality of your program
 - Only large leaks can be found by monitoring the memory footprint
 - Even if you know a leak is exists it is very difficult to determine which code is the culprit
- This is such an issue that many languages such as java will not allow the programmer to control memory allocation and deallocation

Memory Leak Strategy

- Besides being a perfect programmer there are tools to help solve this memory leak issue
- For Visual Studio there is a memory leak detection tool
 - Requires an additional define
 - Use through the debugger tool
 - A howTo for leak detection is uploaded to this lecture set
 - The tool is very limited compared to Valgrind
- For g++ based compilers (linux, mac)
 - Valgrind